

# A Short Introduction to CARTs

Andreas Kagoshima, *M. Sc. Student, Karlsruhe Institute of Technology KIT*

First of all some important definitions:




Figure 1: CART: Hastie, Tibshirani, Friedman- The Elements of Statistical Learning

## CARTs (Classification and Regression Trees)

- They are called **CARTs**, because they can be used for either: Classification and Regression.
- We will focus on **regression trees**.
- Regression trees partition the predictor space into simple regions and do a constant linear regression for each region. (Figure 1: Top left: non-admissible regions, top right: admissible regions).
- Since the set of splitting rules used to partition the predictor space can be summarized in a tree, these types of approaches are known as decision-tree methods.
- Tree-based methods are simple and useful for interpretation.
- However they typically are not competitive with the best supervised learning approaches in terms of prediction accuracy.
- Hence we also discuss the **ensemble methods**: **Bagging**, **random forests**, and **boosting**.
- Given a partition  $R_1, \dots, R_J$  and data-points  $[x^{(i)}, y^{(i)}]^T \in \mathbb{R}^{(p+1)}$ . Then the resulting (fitted) regression model is given by

$$f(x) = \sum_{j=1}^J \hat{c}_j \mathbb{1}_{\{x \in R_j\}},$$

where the  $\hat{c}_j$  are given by

$$\hat{c}_j = \text{ave}(y^{(i)} \mid x^{(i)} \in R_j) = \frac{\sum_{i=1}^N y^{(i)} \mathbb{1}_{x^{(i)} \in R_j}}{\sum_{i=1}^N \mathbb{1}_{x^{(i)} \in R_j}}.$$

This is the optimal choice under quadratic loss (quadratic scoring function).

### Tree-building process

- To find the best partition  $R_1, \dots, R_J$ , under quadratic loss

$$\sum_{j=1}^J \sum_{\{i: x^{(i)} \in R_j\}} (y^{(i)} - \hat{c}_j)^2$$

is computationally infeasible, because one would have to consider every possible partition of the feature space into  $J$  boxes.

- For this reason, one often uses the **top-down greedy approach** that is also known as **recursive binary splitting**.
- The approach is called top-down because it begins at the top of the tree and then successively splits the predictor space; each split is indicated via two new branches further down on the tree.
- It is greedy, because at each step of the tree-building process, the best split is made at the particular step, rather than looking ahead and picking a split that will lead to a better tree in some future step.
- Formally this algorithm can be described as follows: Let  $x_j$  (for some  $j \in \{1, \dots, p\}$ ) be the  $j$ -th row of the vector  $x \in \mathbb{R}^p$  in the prediction space ( $x_j =$  **splitting variable**) and let  $s \in \mathbb{R}$  be an one-dimensional bound ( $s =$  **split point**). We can then define a pair of half-planes by the splitting rule

$$R_1(j, s) = \{x \in \text{prediction space} \mid x_j \leq s\}, \quad R_2(j, s) = \{x \in \text{prediction space} \mid x_j > s\}.$$

We then seek the splitting variable  $j^*$  and the split point  $s^*$ , so that

$$\begin{aligned} (j^*, s^*) &= \underset{(j,s)}{\operatorname{argmin}} \left[ \min_{c_1} \sum_{\{i: x^{(i)} \in R_1(j,s)\}} (y^{(i)} - c_1)^2 + \min_{c_2} \sum_{\{i: x^{(i)} \in R_2(j,s)\}} (y^{(i)} - c_2)^2 \right] \\ &= \underset{(j,s)}{\operatorname{argmin}} \left[ \sum_{\{i: x^{(i)} \in R_1(j,s)\}} (y^{(i)} - \hat{c}_1(j,s))^2 + \sum_{\{i: x^{(i)} \in R_2(j,s)\}} (y^{(i)} - \hat{c}_2(j,s))^2 \right]. \end{aligned}$$

The inner minimization problems are always solved by

$$\hat{c}_1(j, s) = \text{ave}(y^{(i)} \mid x^{(i)} \in R_1(j, s)), \quad \hat{c}_2(j, s) = \text{ave}(y^{(i)} \mid x^{(i)} \in R_2(j, s)).$$

- For each splitting variable, the determination of the split point  $s$  can be done very quickly and hence by scanning through all of the splitting variables  $x_j$ , determination of the best pair  $(j, s)$  is feasible.
- One usually uses binary splitting, stopping only, when each terminal node has fewer than some minimum number of observations.

### Pruning large trees

- The process above may produce good predictions on the training data, but is likely to overfit the data, leading to poor test set performance. Why?
- A smaller tree with fewer splits (that is, fewer regions  $R_1, \dots, R_J$ ) might lead to lower variance and better interpretation at the cost of a little bias.

- One possible alternative to the process described above is to grow the tree only so long as the decrease in the RSS due to each split exceeds some (high) threshold.
- This strategy will result in smaller trees, but is too short-sighted: A seemingly worthless split early on in the tree might be followed by a very good split- that is, a split that leads to a large reduction in RSS later on.
- A better strategy is to grow a very large tree  $T_0$  and then prune it back in order to obtain a subtree.
- Cost complexity pruning- also known as weakest link pruning- is used to do this.
- We consider a sequence of trees indexed by a non-negative tuning parameter  $\alpha$ . For each value of  $\alpha$  there corresponds a subtree  $T \subset T_0$  such that


$$\sum_{m=1}^{|T|} \sum_{\{i: x^{(i)} \in R_m\}} (y^{(i)} - \hat{c}_m) + \alpha|T|$$

Here  $|T|$  indicates the number of terminal nodes of the tree  $T$ ,  $R_m$  is the rectangle (i.e. the subset of predictor space) corresponding to the  $m$ -th terminal node, and  $\hat{c}_m$  is the mean of the training observations in  $R_m$ .

- The tuning parameter  $\alpha$  controls a trade-off between the subtree's complexity and its fit to the training data.
- We select an optimal value  $\hat{\alpha}$  using cross-validation.
- We then return to the full data set and obtain the subtree corresponding to  $\hat{\alpha}$ .

## Bagging

- **Bagging** = Bootstrap aggregation (we use trees in an ensemble).
- Usually the bootstrap is used to estimate the standard error or the bias of an estimator.
- **Bagging** is a general purpose procedure for reducing the variance of a statistical learning method; we introduce it here, because it is particularly useful and frequently used in the context of **decision trees**.
- Recall that given a set of  $n$  independent observations  $Z_1, \dots, Z_n$ , each with variance  $\sigma^2$ , the variance of the mean  $\bar{Z}$  of the observations is given by  $\sigma^2/n$ .
- In other words, averaging a set of observations reduces the variance. Of course, this is not practical because we generally do not have access to multiple training sets.
- Instead, we create **bootstrap data-sets**. (This means we create data-sets of the same size, consisting out of uniformly drawn data-points of the original data-set).



Usually we create a few hundred **bootstrap data-sets**.

- On each **bootstrap data-set** we now create a big tree without pruning. Hence each tree has high variance, but low bias.
- Averaging these  $B$  trees reduces the variance. This is an **ensemble method**.
- In formulas: We generate  $B$  different **bootstrapped data-sets**. We then train our method on the  $b$ -th bootstrapped training set in order to get  $\hat{f}^{*(b)}(x)$ , the prediction at a point  $x$ . We then average all the predictions to obtain

$$\hat{f}_{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*(b)}(x).$$

### Out-of-Bag Error Estimation

- It turns out that there is a very straightforward way to estimate the test error of a bagged model.
- Recall that the key to bagging is that trees are repeatedly fit to bootstrapped sub-sets of the observations. One can show that on average, each bagged tree makes use of around two-thirds of the observations.
- The remaining one-third of the observations not used to fit a given bagged tree are referred to as the **out-of-bag** (OOB) observation.
- We can predict the response for the  $i$ -th observation using each of the trees in which that observation was OOB. This will yield around  $B/3$  predictions for the  $i$ -th observation, which we average.
- The estimate is essentially the LOO cross-validation error for bagging, if  $B$  is large.

### Random Forests (RFs)

- **Random forests** is also an **ensemble method**. **RFs** provide an improvement over bagged trees by way of a small tweak that decorrelates the trees. This reduces the variance when we average the trees.
- As in bagging, we build a number of decision trees on **bootstrapped data-set**.
- But when building these decision trees, each time a split in a tree is considered, a random selection of  $m$  predictors is chosen as split candidates from the full set of  $p$  predictors. The split is allowed to use only one of those  $m$  predictors.
- A fresh selection of  $m$  predictors is taken at each split, and typically we choose  $m \approx \sqrt{p}$ —that is, the number of predictors considered at each split is approximately equal to the square root of the total number of predictors.

### Boosting

- **Boosting** is an **ensemble method**. Like **bagging**, **boosting** is a general approach that can be applied to many statistical learning methods for regression or classification. We only discuss boosting for decision trees.
- Recall that bagging involves creating **bootstrapped data-sets**. Then a separate decision tree is fit to each copy and then combining all of the trees in order to create a single predictive model.
- Notably, each tree is built on a **bootstrap data-set**, independent of the other trees.
- Boosting works in a similar way, except that the trees are grown sequentially: Each tree is grown using information from previously grown trees.
- Boosting algorithm for regression trees

1. Set  $\hat{f}(x) = 0$  and  $r^{(i)} = y^{(i)}$  for all  $i$  in the training set.

2. For  $b = 1, 2, \dots, B$ , repeat:
  - (a) Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d+1$  terminal nodes) to the training data  $(X, r)$ .
  - (b) Update  $f$  by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x).$$

- (c) Update the residuals

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x).$$

3. Output the boosted model

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x).$$

- This is also an **ensemble method**.

### What is the idea behind this procedure?

- Unlike fitting a single large decision tree to the data, which amounts to **fitting the data hard** and potentially overfitting, the boosting approach instead learns slowly.
- Given the current model, we fit a decision tree to the residuals from the model. We then add this new decision tree into the fitted function in order to update residuals.
- Each of these trees can be rather small, with just a few terminal nodes, determined by the parameter  $d$  in the algorithm.
- By fitting small trees to the residuals, we slowly improve  $\hat{f}$  in areas where it does not perform well. The shrinkage parameter  $\lambda$  shows the process down even further, allowing more and different shaped trees to attack the residuals.

### Tuning parameters for boosting

- The number of trees  $B$ : Unlike **bagging** and **random forests (RF)**, boosting can overfit if  $B$  is too large, although this overfitting tends to occur slowly if at all. We use cross-validation to select  $B$ .
- The shrinkage parameter  $\lambda$  is a small positive number. This controls the rate at which boosting learns. Typical values are 0.01 or 0.001, and the right choice can depend on the problem. Very small  $\lambda$  can require using a very large value of  $B$  in order to achieve good performance.
- The number of splits  $d$  in each tree, which controls the complexity of the boosted ensemble. Often  $d = 1$  works well, in which case each tree is a stump, consisting of a single split and resulting in an additive model. More generally  $d$  is the interaction depth, and controls the interaction order of the boosted model, since  $d$  splits can involve at most  $d$  variables.

### CARTs and RFs for Density Forecasting

- First of all Regression trees and Random Forests are typically used for regression.
- But in return they can also be used for density forecasting.
- Consider a tuple  $[x, y]^T \in \mathbb{R}^{p+1}$  and consider the case of a regression tree. Then

$$P[y = y^{(k)} | x \in \text{node } t] \quad (k \hat{=} \text{"class" } k)$$

can be estimated with the estimator

$$\hat{f}(y = y^{(k)} | x \in \text{node } t) := \frac{\#\{\text{data-points } y : y = y^{(k)} \wedge x \in \text{node } t\}}{\#\{\text{data-points } y : x \in \text{node } t\}}.$$

- Consider a tuple  $[x, y]^T \in \mathbb{R}^{p+1}$ ,  $B$  bootstrap data-sets and consider the case of a RF. Then

$$P[y = y^{(k)} | x \in \text{node } t] \quad (k \hat{=} \text{"class" } k)$$

can be estimated with the estimator (created from the  $b$ -th bootstrap data-set)

$$\hat{f}^{(b)}(y = y^{(k)} | x \in \text{node } t) := \frac{\#\{\text{data-points } y \in b\text{-th bootstrap data-set} : y = y^{(k)} \wedge x \in \text{node } t\}}{\#\{\text{data-points } y \in b\text{-th bootstrap data-set} : x \in \text{node } t\}},$$

and through an ensemble approach one could use

$$\hat{f}_{\text{RF}}(y = y^{(k)} | x \in \text{node } t) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{(b)}(y = y^{(k)} | x \in \text{node } t)$$

as a density forecast.

## Literature

- The Elements of Statistical Learning- Hastie, Tibshirani, Friedman,
- An Introduction to Statistical Learning- James, Witten, Hastie, Tibshirani,
- Lecture Notes: "Statistisches Lernen"- Universität Freiburg, 2014,
- Stanford University Online Lecture: "Statistical Learning".